

# CSCE 3402 Spring 2021

## Exercise 10 – Parallel Fork

### Redundant Fork for Fault Tolerance

Assigned: Tuesday, April 20<sup>th</sup> in Lab

*Due: Tuesday, May 18<sup>th</sup>*

No Delayed Submission is Allowed.

#### Goals

In this lab exercise you are required to modify the Linux kernel. This lab exercise is an individual lab exercise that you should work on it on your own and is designed to enrich your kernel development skills. In this lab exercise you will need to read parts for the linux kernel code, understand it, and modify/add to it to achieve the target goal. In this lab exercise you will add a number of system calls to achieve some fault tolerance mechanism that will allow a killed process to resume execution under certain conditions; simply you will be implementing a new version of the fork system call to provide redundancy and fault tolerance.

#### Basics

In this lab exercise you will have to add a number of new system calls; one main important system call and a couple of another two helper system calls. The main system call you need to add is the **pfork** system call. The **pfork** system call is like the fork system call, but it is not :)

The **pfork** system call should fork two children processes instead of one. One child process should be in a runnable state and on the run queue, while the other one should be in a suspended state on a wait queue. From now on, we will refer to the runnable one as active and the other suspended one as standby. As long as the active process is running the standby will stay in the suspended state. As soon as the active one is terminated the standby should start executing and should become active.

All your code should be submitted in the form of one or more patches. You should learn how to create a patch for multiple recursive directories. This will entail taking a copy of the directories you are going to work on before starting to apply your amendments and use the **diff** command to create the patch(es). It is worthy reading the man page of the diff command.

## Details

Your kernel amendments should achieve the following as a minimum:

1. You should extend the process block data structure to contain extra fields as a minimum (you might need to add more fields to achieve your target):
  1. **pfork\_standby\_pid**: this field should be set in the active process to contain the pid of the standby one, and it should only be set to a value other than zero if **pfork** is used.
  2. **pfork\_active\_pid**: this field should be set in the standby process to contain the pid of the active one, and it should only be set to a value other than zero if **pfork** is used.
  3. **pfork\_status**: a status value that can be set through a system call.
2. Add a new system call and name it **pfork**.
3. **pfork** should perform the following:
  1. Fork two child processes instead of one.
  2. Put the first one on the active run queue.
  3. Put the second one on the wait queue and store in its `pfork_active_pid` the pid of the active child.
4. Add the following system calls:
  1. **set\_pfork\_status**: sets the value of the `pfork_status` to some value based on the program logic. When invoked from the active process it will set the `pfork_status` of both the active and standby processes.
  2. **get\_pfork\_status**: returns the value of the `pfork_status`.
  3. **get\_pfork\_sibling\_pid**: return `pfork_standby_pid` when invoked by the active process and `pfork_active_pid` when invoked by the standby process.
  4. **pfork\_who**: returns 1 when invoked from the active process, 2 when invoked from the standby, and 0 when invoked from a process that was not created by the `pfork` system call.
5. There are set of **SYSCALL\_DEFINE** macros that you should learn and read about to be able to add new systems calls with the right number of arguments.
6. You will also need to learn about some kernel APIs such as **find\_get\_task\_by\_vpid**, **task\_tgid\_vnr**, **do\_fork**, and more.
7. When the active child terminats for any reason the standby process should be unsuspended and start execution; based on the code of the process and how it utilize the **pfork\_status** execution should be resumed, Please refer to the example code presented at the end of the lab exercise document.
8. You will essentially need to hook the exit system call and any termination signals that are applicable.
9. You are required to build a user side library and call it **pfork** (**pfork.h** and **pfork.c**) which contains all the wrapping functions for the above required system calls. The function wrappers should utilize the **syscall** library function to invoke kernel level system calls from the user space.

10. Some of the main files that you should read, understand, and modify are ***core.c***, ***sched.h***, ***fork.c***, and more.
11. You will need to change a number of make files as well as creating some new make files.

Make sure you setup your development environment as follows:

1. First of all you will need to get the kernel source. For consistency we will avail for you the kernel source version that you should work on so we have a common version for the kernel source that will be used by all the students. Please download it from here:  
<https://drive.google.com/drive/folders/1UxqWbnf4FO9a4P85Xn9idhfb8kgzEYFB?usp=sharing>
2. Moreover you should build a virtual machine using VirtualBox and install on it the latest linux distribution that you prefer, I use mint; **this is not a must.**
3. Make sure you install all needed packages for kernel development and update your environment to the latest using apt-get or similar package management tools; you should be able to dig for that.
4. Copy the downloaded kernel and save it inside your virtual machine.
5. You need to go through one round of compiling the downloaded kernel successfully and booting from it before you start working and modifying code.
6. Make sure that you make a copy of each directory you modify files in so you can generate a patch using **diff**.

## Deliverables

Everything should be submitted to Blackboard. Include a Readme file to explain anything unusual to the TA. Your kernel code must be included in one or more patch files. Other associated files, such as the user space library **pfork** code should be included in .c and .h files and stored under a separate directory. A test program that utilizes the **pfork** system call and the other required system calls which demonstrate the required functionalities should also be included. Create a single directory that include your patches, .h and .c files, readme file, and design document. Finally try to organize your files under different directories based on their types.

Your design document should be called **design.pdf** and it should be in PDF format and should be in the same directory. Formats other than PDF are not acceptable; please convert other formats (Word, LaTeX, HTML, ...) to PDF. Your design document should describe the design of your lab exercise in enough details that a knowledgeable programmer could duplicate your work. This includes description of the data structures you use, all non-trivial algorithms and formulas, and a description of each function including its purpose, inputs, outputs, and assumptions you made about the inputs or outputs.

### **Deliverables Summary:**

1. Design Document ( .pdf)
2. README.txt
3. Patch files.

### **How to submit:**

Compress all your work: source code, report, readme file, and any extra information into a zip archive. You should name your archive in the specific format <Student\_ID>\_<Name>\_Lab10.zip. Finally, upload your code to blackboard.

### **Grade**

This lab exercise is worth 20% of the overall course grade. The lab exercise will be graded on a 100% grade scale, and then will be scaled down to the 10% its worth. The grading of the lab exercise will be broken down as follows:

1. 60 % for the correctness of the functionality and the quality of your code.
2. 15 % for the quality of your inline documentation and the readme file.
3. 25 % for the design document.

### **Important Note**

One of the main objectives of this lab exercise, as stated above, is to learn how to hack the kernel code, and this requires you to read the kernel code before modification. To set the right expectations, getting help in how to navigate the kernel code and understand the needed parts for this lab exercise will defeat the whole purpose of the lab exercise. Consequently, you should expect less help from the TA regarding that. Simply you need to read the code and try your best to understand it before coming to us for questions. We will be happy to help you with your questions, but you need to show that you did considerable amount of effort. **Also note that compiling the kernel is a time-consuming process and you need to start right away and do not waste any time.**

## Sample User Space Program

```
#include <pfork.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/wait.h>

int main ( int argc, char ** argv )
{
    pid_t cpid = pfork();
    if (cpid !=0 )
        for (;wait(NULL) > 0 ;);
    else
    {
        char who[10];
        memset (who,10,0);
        if ( pfork_who() == 2 )
            strcpy (who,"STANDBY");
        else strcpy (who,"ACTIVE");
        for ( int i = 0 ; i < 5 ; i ++ )
        {
            long int status = get_pfork_status();
            if ( status == i )
            {
                printf ("%s: Current Status is %d\n",who,status);
                sleep ((i+1)*10);
                set_pfork_status(i+1);
                printf ("%s: Set Status is %d\n",who,status+1);
            }
        }
    }
    return 0;
}
```